


The hackathon data analysis is a summary of the EDA activities you have completed. This should include data cleaning, pre-processing, exploratory analysis, and insights you have generated from your data. Each group will submit and demonstrate one notebook (in ipynb and pdf) that comprises the following sections:

- 1.Domain: Explain the domain/subdomain of this project (use images)
  - 2.Data: Go over 1 row of the data explaining relevant columns
  - 3.Insights: Show intuitive and non-intuitive insights generated from your EDA
  - 4.Expected Objective: Sample input and output
- 

## Domain

### Cloud Computing

**Cloud Computing (NIST):** a model for enabling on-demand network access to a shared pool of computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction.



### 5 Key Characteristics of Cloud Computing

1. Broad Network Access
2. Measured Service
3. On Demand Self Service

4. Shared Resource Pooling
5. Rapid Elasticity

### 3 Service Models

1. **Software as a Service (SaaS)**: Provider's applications running on a cloud infrastructure
2. **Platform as a Service (PaaS)**: Consumer-created or acquired applications created using tools supported by the provider
3. **Infrastructure as a Service (IaaS)**: Consumer is able to deploy and run arbitrary software on storage, networks, etc.

## Leading Cloud Service Providers

### 1. Amazon Web Services (AWS)



### 1. Microsoft Azure



## 1. Google Cloud Platform



**Cloud Migration:** the process of moving a company's digital assets, services, databases, IT resources, and applications either partially, or wholly, into the cloud.

## Cloud Management



Cloud Deployment



Cloud Security



Network and Storage



Cloud Monitoring and Reporting



Backup and Disaster Recovery



Infrastructure Set Up



Business Continuity



Cloud Migration Services

**MultiCloudX:** Austin based start-up that functions as a third-party cloud manager to manage cloud storage, operations, and costs



Unleash your business potential  
with the cloud

# Data

```
In [1]: import numpy as np
import pandas as pd

import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: pd.set_option('display.max_columns', None)
```

```
In [3]: dfCCO = pd.read_csv(r".\Data\cco_cost_monthly.csv")
```

```
In [4]: dfPYCO = pd.read_csv(r".\Data\pyco_cost_monthly.csv", low_memory=False)
```

```
In [5]: dfXCO = pd.read_csv(r".\Data\xco_cost_monthly.csv")
```

```
In [6]: dfCCO.head()
```

```
Out[6]:
```

	cloud_resource_id	usage_account_id	billing_account_id	provider_code	usage_amount	currency_code	tax	sub_total	total_cost	system_cur
0	NaN	751355800400	751355800400	aws	1.0	USD	0.0	0.0	0.03	
1	NaN	751355800400	751355800400	aws	2.0	USD	0.0	0.0	0.07	
2	NaN	751355800400	751355800400	aws	1.0	USD	0.0	0.0	0.00	
3	arn:aws:kms:us-east-1:751355800400:key/5fedb8a...	751355800400	751355800400	aws	18.0	USD	0.0	0.0	0.00	
4	arn:aws:cognito-idp:us-east-1:751355800400:use...	751355800400	751355800400	aws	2.0	USD	0.0	0.0	0.00	

cloud\_resource\_id usage\_account\_id billing\_account\_id provider\_code usage\_amount currency\_code tax sub\_total total\_cost system\_currency\_code

Useful: total\_cost, service\_name Maybe: usage\_amount, product\_name, usage\_type, location\_id, cost\_type, instance\_type Not Useful: cloud\_resource\_id, usage\_account\_id, billing\_account\_id, provider\_code, currency\_code, tax, sub\_total, system\_currency\_code, conversion\_rate, converted\_total\_cost (redundant), service\_code, product\_sku, availability\_zone, location\_id, usage\_type\_group,

# Insights

---

## EDA

Now that the functions that we will use have been created, let's look at some key features of our data. First, we combine the dataframes to get general ideas about the full data we were provided.

```
In [7]: dfTotal = pd.concat([dfCCO, dfPYCO, dfXCO], axis=0)
dfTotal.shape
```

```
Out[7]: (984010, 44)
```

So, almost 1 million rows and 44 columns! We will need to aggregate this data further to get a better sense of what it means. Let's now look at the general variability of our columns by seeing how many unique values each has.

```
In [8]: dfTotal.nunique()
```

```
Out[8]: cloud_resource_id      279088
usage_account_id              4
billing_account_id            3
provider_code                  1
usage_amount                  111997
currency_code                  1
tax                            1
sub_total                      1
total_cost                     132600
system_currency_code           1
conversion_rate                 1
converted_total_cost           132600
product_name                    47
service_code                     41
service_name                     40
product_sku                     598
availability_zone               13
location_id                     18
usage_type                      525
usage_type_group                61
cost_type                       6
```



```
instance_type          52
blended_rate           479
unblended_rate         143
unblended_cost        127870
category               60
clock_speed            11
from_location          24
to_location            28
transfer_type          6
pricing_rate_id        2866
public_ondemand_cost   66222
public_ondemand_rate    138
pricing_term           2
pricing_unit           56
memory                17
operating_system       3
vpcu                   6
physical_processor     11
volume_type            10
storage                6
storage_class          3
storage_media          6
invoice_month          19
dtvpe: int64
```


---

Since features like "conversion\_rate" and "provider\_code" are unvaried, they can be discarded as features as invariate features are worthless and can be dropped. Next, let's look at a quick correlation matrix of our full dataset to see if anything has a high level of correlation with the total\_cost that isn't fully redundant.

```
In [9]: def plot_corr(dataframe, size=10):
        """
        Plots a correlation matrix as a heat map
        """
        corr = dataframe.corr()
        fig, ax = plt.subplots(figsize=(size, size))
        im = ax.matshow(corr, vmin = -1.0, vmax = 1.0, cmap = "bwr")
        plt.xticks(range(len(corr.columns)), corr.columns, rotation = 90);
        plt.yticks(range(len(corr.columns)), corr.columns);
        plt.colorbar(im, orientation = 'vertical')
        plt.title('Correlation Matrix')
```

```
In [10]: plot_corr(dfTotal)
```

Correlation Matrix



It looks like there's some light correlation to usage and the vcpu count... nothing too useful, although we may need to make another correlation matrix after some data augmentation. Next, let's look at the %NaN for each column to see which columns are generally unused in the dataset.

```
In [11]: dfTotal.isna().sum() / dfTotal.shape[0]
```

```
Out[11]: cloud_resource_id    0.003146
usage_account_id             0.000000
billing_account_id           0.000000
```

```

provider_code      0.000000
usage_amount       0.000000
currency_code      0.000000
tax                0.000000
sub_total          0.000000
total_cost         0.000000
system_currency_code 0.000000
conversion_rate    0.000000
converted_total_cost 0.000000
product_name       0.000000
service_code       0.084868
service_name       0.084906
product_sku        0.000605
availability_zone   0.995690
location_id        0.000605
usage_type         0.084868
usage_type_group   0.989718
cost_type          0.000000
instance_type      0.996030
blended_rate       0.000000
unblended_rate     0.000000
unblended_cost     0.000000
category           0.000000
clock_speed        0.996988
from_location      0.577076
to_location        0.577076
transfer_type      0.577076
pricing_rate_id    0.000605
public_ondemand_cost 0.000000
public_ondemand_rate 0.000000
pricing_term       0.097949
pricing_unit       0.000605
memory            0.996018
operating_system   0.996992
vpcu               0.996012
physical_processor 0.996180
volume_type        0.996372
storage            0.996179
storage_class      0.998962
storage_media      0.985840
invoice_month      0.000000
dtype: float64

```

It looks like the 3 columns for transfers all need one another to exist. Then, the columns like "instance\_type", "volume\_type", and "operating\_system" all likely point towards the amount of line items that are actual VMs that have been spun up rather than associated costs or abstracted cloud services. It could perhaps be possible to use the amount of rows we have for each of those columns as advisors of fixed costs as certain server instances such as domain controllers, FTP servers, and application servers require high availability.

```
In [12]: dfPYCO.sort_values("total_cost", ascending=False)[["invoice_month", "product_name", "total_cost"]]
```

```
Out[12]:
```

invoice_month	product_name	total_cost
---------------	--------------	------------

	invoice_month	product_name	total_cost
<b>65279</b>	2019-09-01	Amazon Elastic Compute Cloud	3284.000000
<b>28934</b>	2019-09-01	AWS Support (Business)	1000.000000
<b>96649</b>	2019-10-01	AWS Support (Business)	1000.000000
<b>114403</b>	2019-10-01	Amazon Elastic File System	879.020880
<b>29088</b>	2019-09-01	Amazon Elastic File System	861.036975
...	...	...	...
<b>567528</b>	2019-12-01	Amazon Elastic Compute Cloud	-3.050000
<b>96842</b>	2019-10-01	Amazon Elastic Compute Cloud	-100.000000
<b>28932</b>	2019-09-01	AWS Support (Business)	-130.774159
<b>96648</b>	2019-10-01	AWS Support (Business)	-708.380000
<b>28935</b>	2019-09-01	AWS Support (Business)	-1000.000000

Here, we see some interesting behaviour where it looks like there were certain charges put forward that then had to be fixed (look at the indices of the largest negative charges and how they're +1 of the large charges).

## Plotting

Now, let's make some figures to get an idea of how the big hitter services are affecting costs month over month.

```
In [13]: def hitter_monthly(df, product):
        """
        Grabs the total cost incurred by the big hitter (product) for each month of the DataFrame, df
        """
        dfFin = df.loc[df["product_name"] == product].groupby("invoice_month").agg({"total_cost": "sum"})
        return dfFin
```

In [39]:

```
def plot_n_hitters(n_hitters):
    """
    Plots the N biggest overall hitters and the aggregate cost
    """
    top_hitters = dfTotal.groupby("product_name").agg({"total_cost": "sum"}).sort_values("total_cost", ascending=False)
    company_dfs = [dfXCO, dfPYCO, dfCCO]
    company_names = ["XCO", "PYCO", "CCO"]
    final_df_list = []
    fig, ax = plt.subplots(3,2, figsize=(14, 8.5))
    for index in range(len(company_names)):
        inter_list = []

        for hitter in range(len(top_hitters)):
            new_monthly_df = hitter_monthly(company_dfs[index], top_hitters[hitter])
            ax[index][0].plot(new_monthly_df)
            ax[index][0].tick_params("x", rotation=90)
            inter_list.append(new_monthly_df)

        ax[index][0].set_title(str(company_names[index]) + " big hitters monthly cost")
        ax[index][0].legend(top_hitters, bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.)
        final_df_list.append(inter_list)
        agg_df = company_dfs[index].groupby("invoice_month").agg({"total_cost": "sum"})
        ax[index][1].plot(agg_df, color="red")
        ax[index][1].set_title(str(company_names[index]) + " total monthly cost")
        ax[index][1].tick_params("x", rotation=90)

    fig.tight_layout()
    fig.savefig(r".\figures\graphs_mod" + str(n_hitters) + ".png")
    return final_df_list
```

```
In [40]: def plot_n_hitters_bycorp(n_hitters):
        """
        Plots the top N big hitters on a company-by-company basis (possibly different for each company)
        """
        company_dfs = [dfXCO, dfPYCO, dfCCO]
        company_names = ["XCO", "PYCO", "CCO"]
        final_df_list = []
        fig, ax = plt.subplots(3,2, figsize=(14, 8.5))
        for index in range(len(company_names)):
            inter_list = []
            top_hitters = company_dfs[index].groupby("product_name").agg({"total_cost": "sum"}).sort_values("total_cost", ascending=False)

            for hitter in range(len(top_hitters)):
                new_monthly_df = hitter_monthly(company_dfs[index], top_hitters[hitter])
                ax[index][0].plot(new_monthly_df)
                ax[index][0].tick_params("x", rotation=90)
                inter_list.append(new_monthly_df)

            ax[index][0].set_title(str(company_names[index]) + " big hitters monthly cost")
            ax[index][0].legend(top_hitters, bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.)
            final_df_list.append(inter_list)
            agg_df = company_dfs[index].groupby("invoice_month").agg({"total_cost": "sum"})
            ax[index][1].plot(agg_df, color="red")
            ax[index][1].set_title(str(company_names[index]) + " total monthly cost")
            ax[index][1].tick_params("x", rotation=90)

        fig.tight_layout()
        fig.savefig(r".\figures\company_hitters_" + str(n_hitters) + ".png")
```

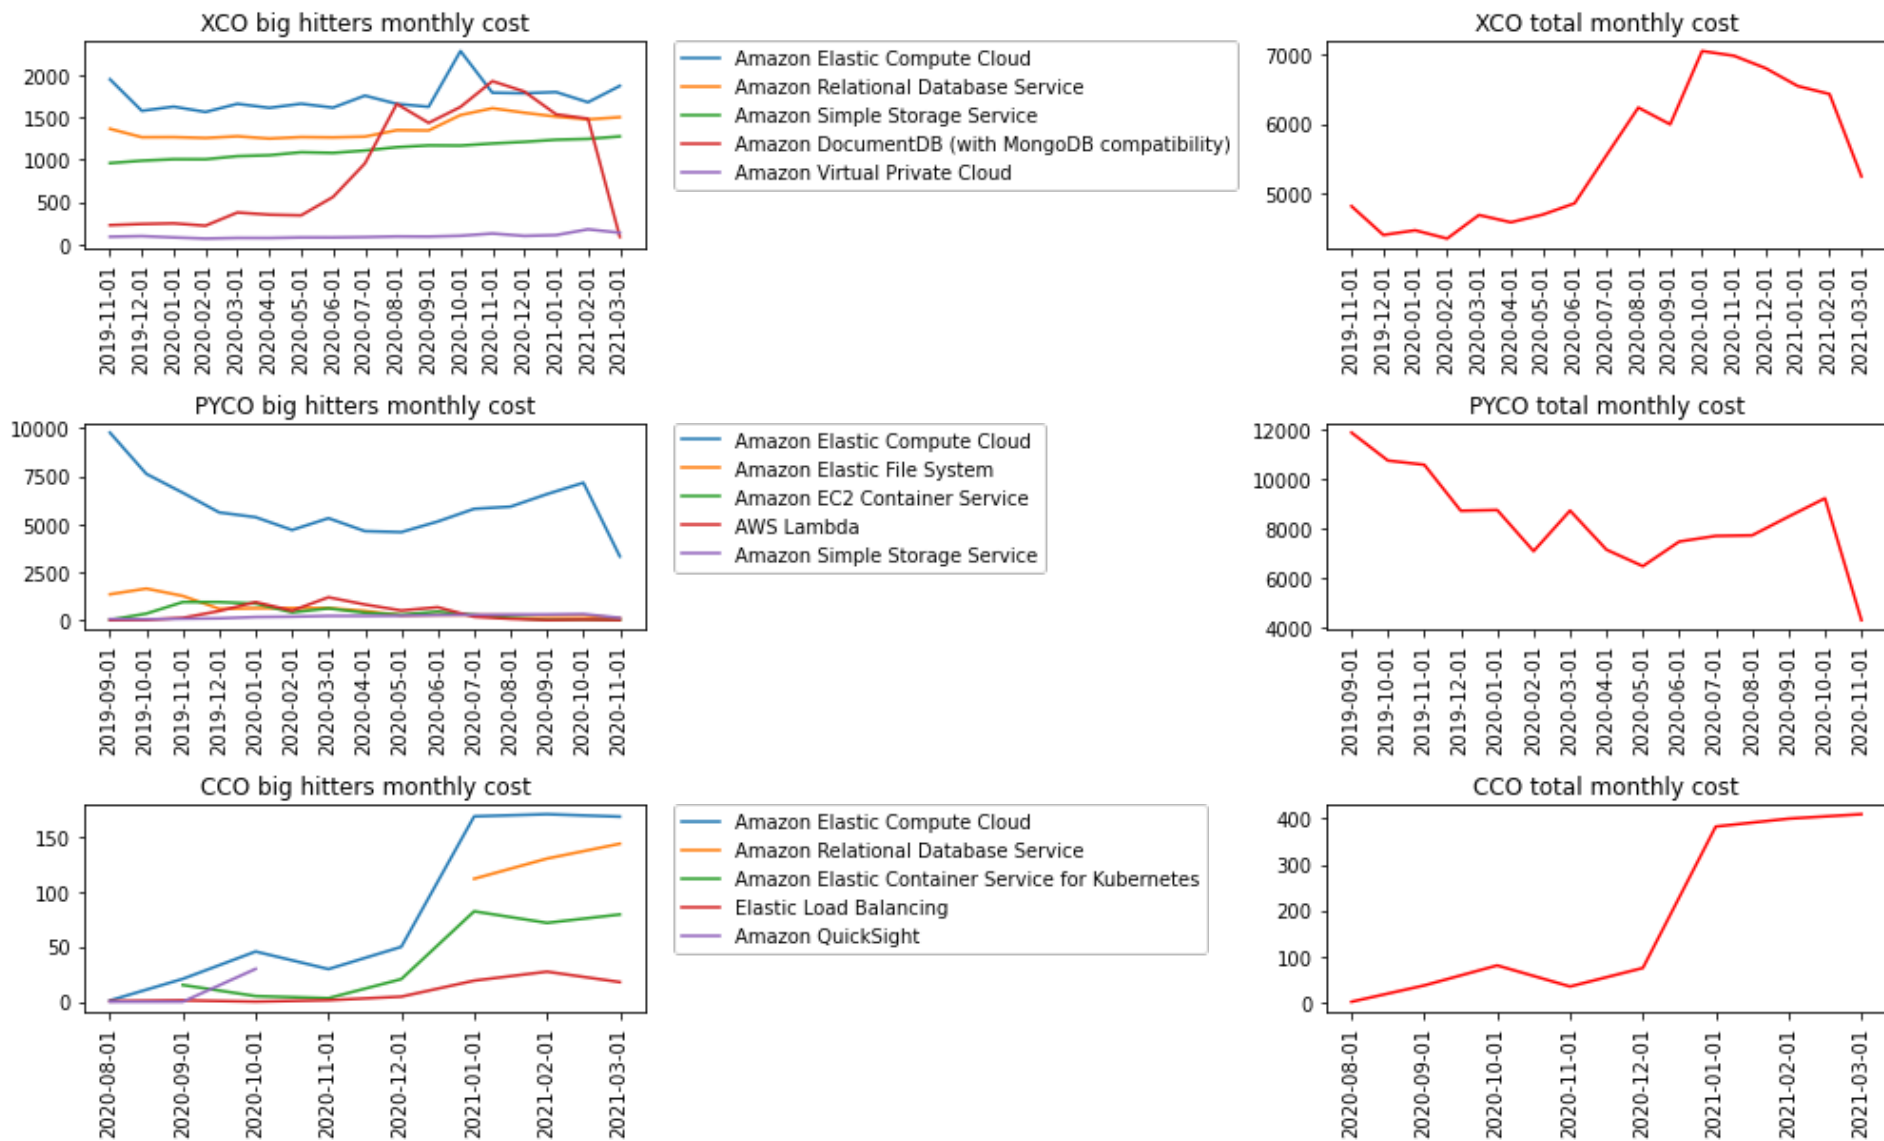
The first plot here is a demonstration of how the top 5 overall "big hitters" affect the aggregate cost (right) for each company. We can see that the top 5 big hitters aren't necessarily the biggest cost drivers for each company. Another interesting insight is how for XCO, the graphical shape seems to be driven more by Amazon DocumentDB even though it's not the largest charge most months, it is the biggest contributor of variance.

```
In [41]: top_5 = plot_n_hitters(5)
```



Here, we look at the top 5 top "big hitters" and see that there were a lot of costs that matter to company with a smaller cloud footprint, like CCO that don't necessarily affect a relatively larger corporation like PYCO.

```
In [42]: plot_n_hitters_bycorp(5)
```



## Data Processing

Now, we will engineer our features to be used in the final model. We are taking the usage and total cost of each product name and taking their sum total for each month. This creates a huge amount of features that can then be later cut down.

Other possible datum:

- # of Windows/Linux/RHEL machines
- Amount of VCPUs/Storage Type/Memory used



```
In [18]: # First start the creation of our final dataframe, starting with total monthly cost
def process_dataframe(df, fillna=True):
    """
    Processes a dataframe from what's initially given to a usable form for our model.
    """
    dfProcessed = df.groupby("invoice_month").agg({"total_cost": "sum"})
    products = dfTotal["product_name"].unique()
    # Loops through all of the unique product names
    for product in products:
        # Grabs the monthly total cost and usage for each product
        df_product_monthly = df.loc[df["product_name"] == product].groupby("invoice_month").agg({"total_cost": "sum", "us
        #Concatenates that onto the total dataframe we have
        dfProcessed = pd.concat([dfProcessed, df_product_monthly], axis=1)
    if fillna==True:
        dfProcessed = dfProcessed.fillna(0)
    return dfProcessed
```

```
In [19]: dfCCO_processed = process_dataframe(dfCCO, fillna=False)
dfCCO_processed
```

```
Out[19]:
```

	total_cost	Elastic Load Balancing cost	Elastic Load Balancing usage	Amazon Elastic Compute Cloud cost	Amazon Elastic Compute Cloud usage	Amazon Simple Storage Service cost	Amazon Simple Storage Service usage	AWS Key Management Service cost	AWS Key Management Service usage	Amazon Cognito cost	Amazon Cognito usage	AmazonC
invoice_month												
2020-08-01	1.717943	0.547634	24.021073	1.147767	24.566741	0.022542	543.011748	0.0	18.0	0.0	2.0	
2020-09-01	37.258579	1.273340	54.133633	20.755028	824.133404	0.021228	2575.908828	0.0	20.0	0.0	1.0	
2020-10-01	80.664613	0.122506	6.000797	45.572186	1189.248625	0.013111	761.534243	0.0	9.0	NaN	NaN	
2020-11-01	34.867281	1.417530	60.004739	29.653838	1292.703062	0.014404	2149.542143	0.0	167.0	0.0	1.0	
2020-12-01	75.113704	4.633152	194.110072	49.838250	1637.684729	0.016703	3426.625776	0.0	233.0	0.0	1.0	
2021-01-01	382.368784	19.185234	803.845366	168.243566	5142.464178	0.025501	5776.833469	0.0	239.0	0.0	57.0	
2021-02-01	399.512312	27.303523	1141.367505	170.194027	5240.705849	0.018575	3482.610250	0.0	269.0	0.0	7.0	
2021-03-01	409.219731	17.845686	747.001553	167.964576	5917.144565	0.019282	3840.988064	0.0	220.0	0.0	9.0	

## CCO

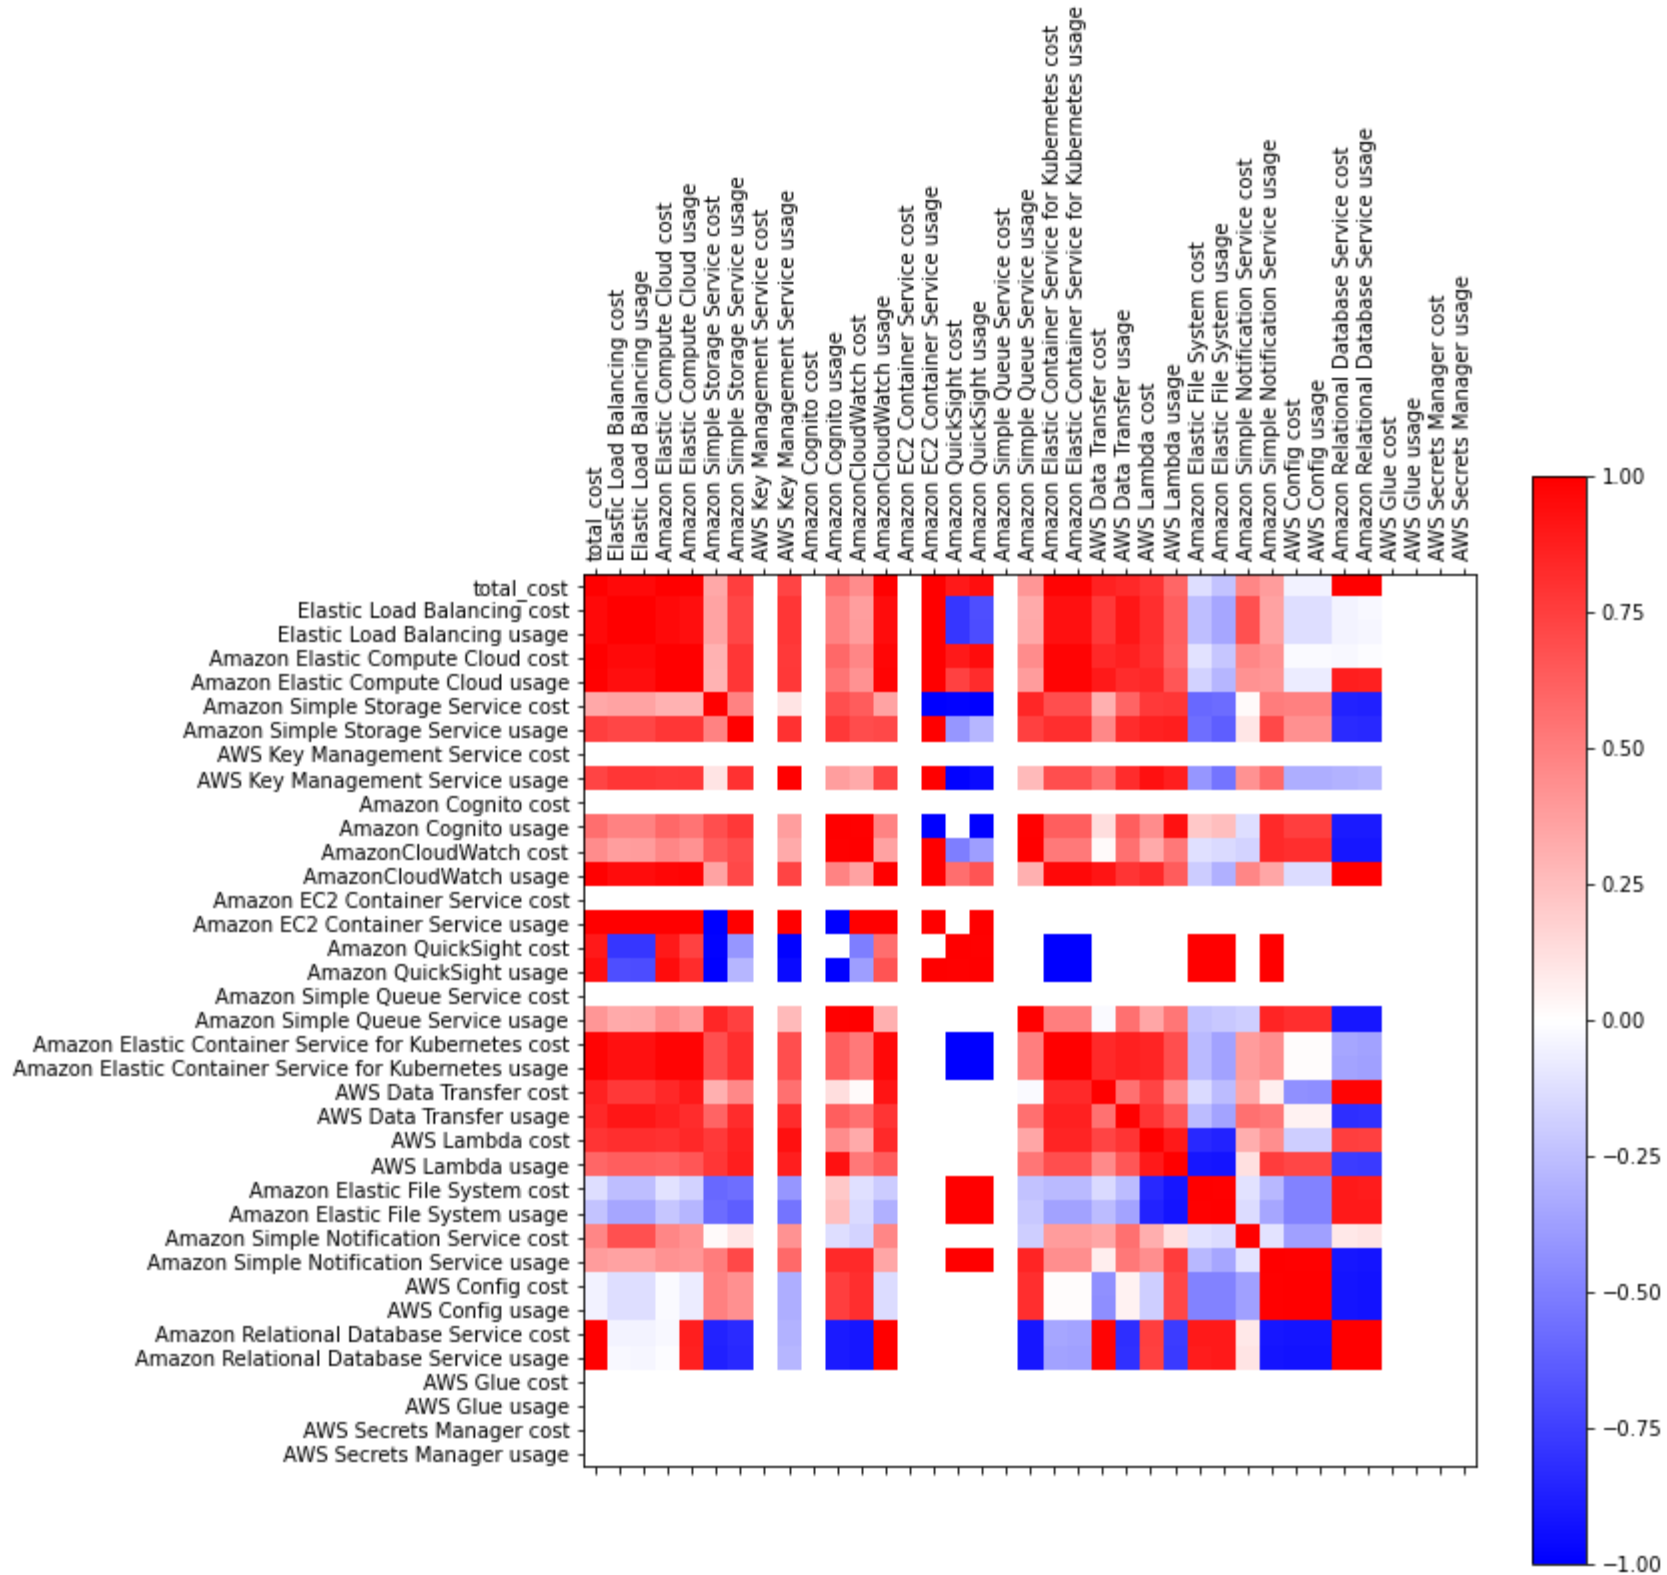
Interesting things to note:

- The only negative correlations lie with the Amazon Elastic File System
- Cost variance seems to be driven strongly by many different services. Likely points to a ramping up of cloud costs, which was reflected in our earlier plots

In [20]:

```
# CCO Correlation Matrix  
plot_corr(dfCCO_processed.dropna(axis=1, how="all"))
```

Correlation Matrix



# PYCO

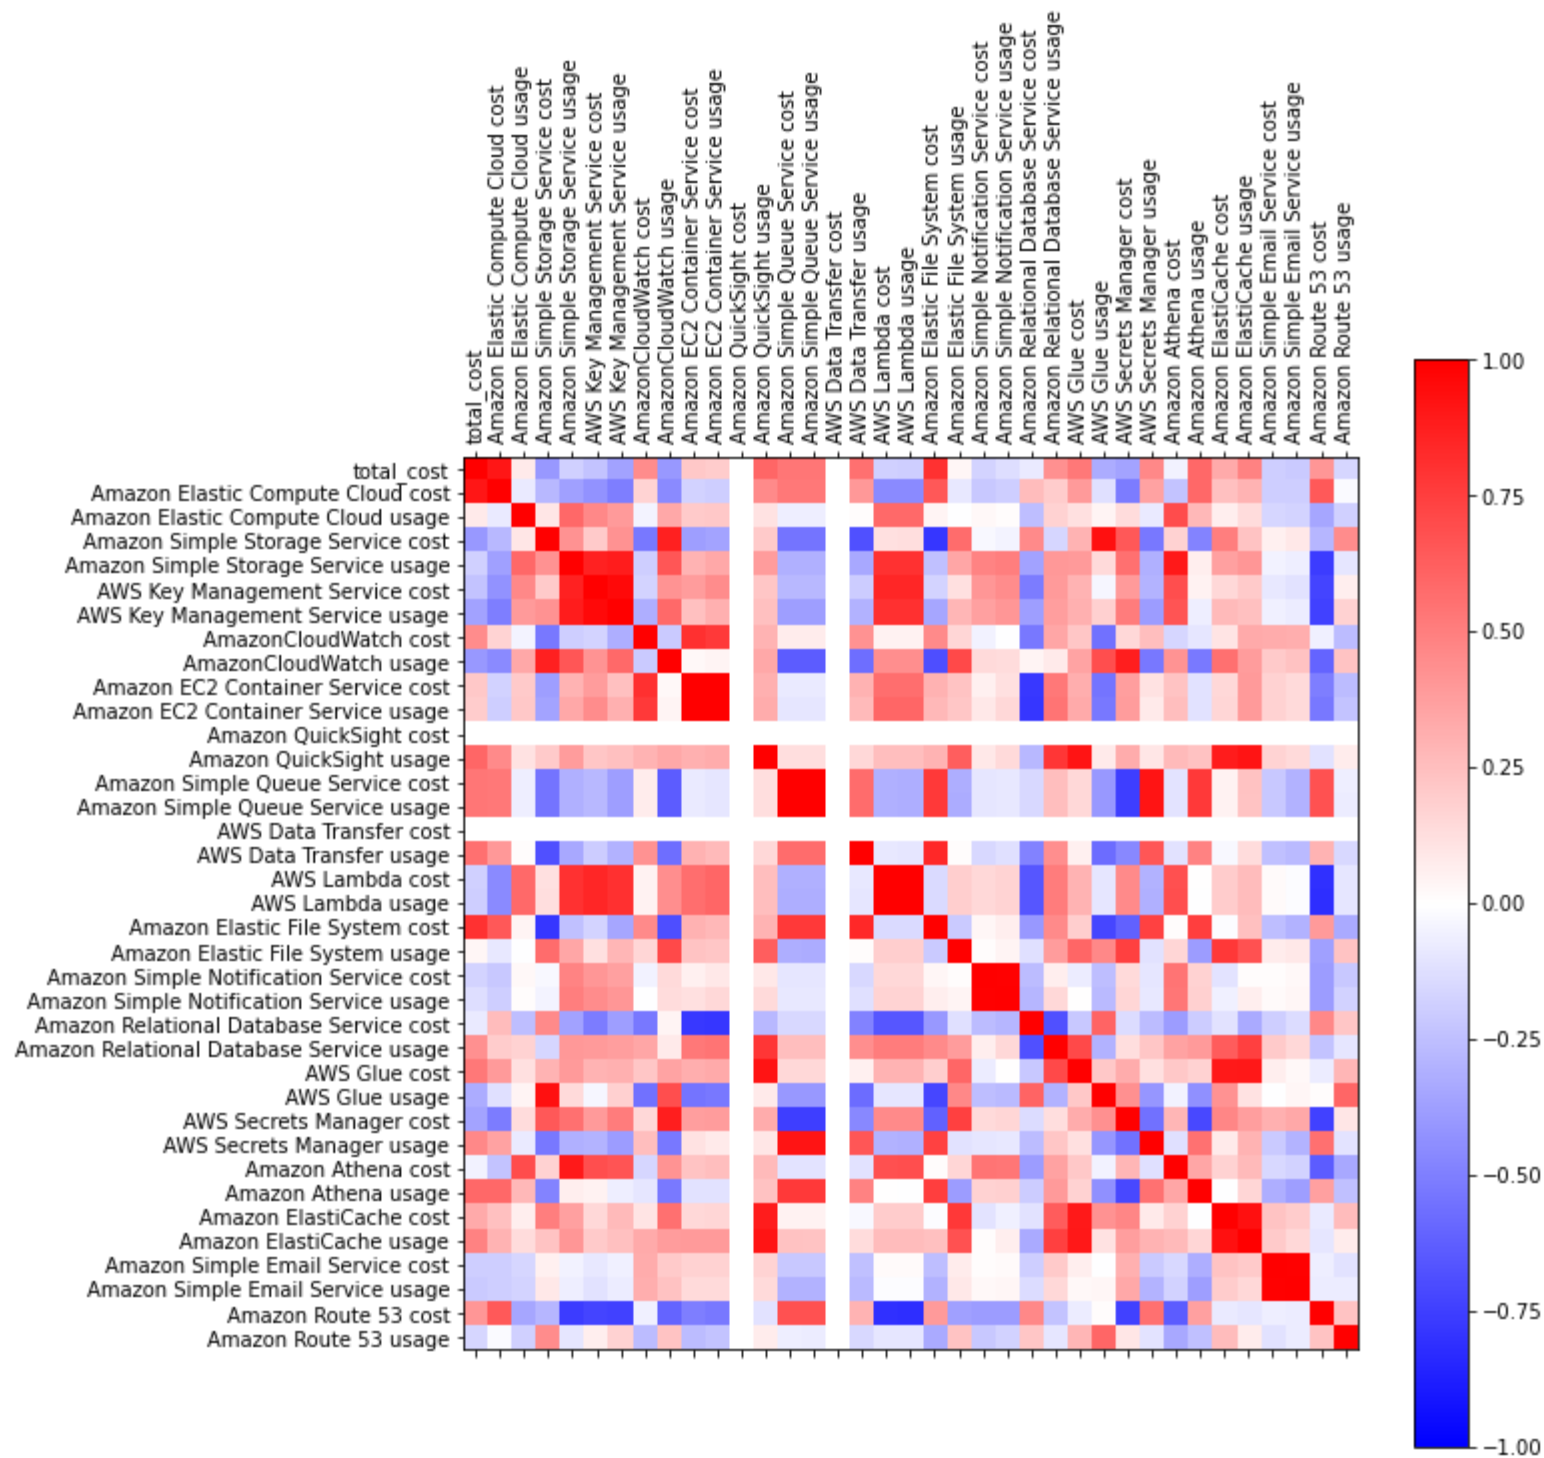
Interesting things to note:

- There's a somewhat large amount of negative correlations for cost, but none are very strong.
- The strongest correlation seems to lie with Amazon Elastic File System costs

In [21]:

```
dfPYCO_processed = process_dataframe(dfPYCO, fillna=False)
# PYCO Correlation Matrix
plot_corr(dfPYCO_processed.dropna(axis=1))
```

Correlation Matrix




Interesting points to note:

- Amazon Lightsail cost has a fairly strong negative correlation with total cost.
- Amazon Relational DB Cost, Amazon simple storage, and Amazon DocumentDB are cost variance drivers. This seems to be a company that intakes a good amount of data and the data intake seems to correlate with increasing costs.

In [22]:

```
dfXCO_processed = process_dataframe(dfXCO, fillna=False)
plot_corr(dfXCO_processed.dropna(axis=1, how="all"))
```

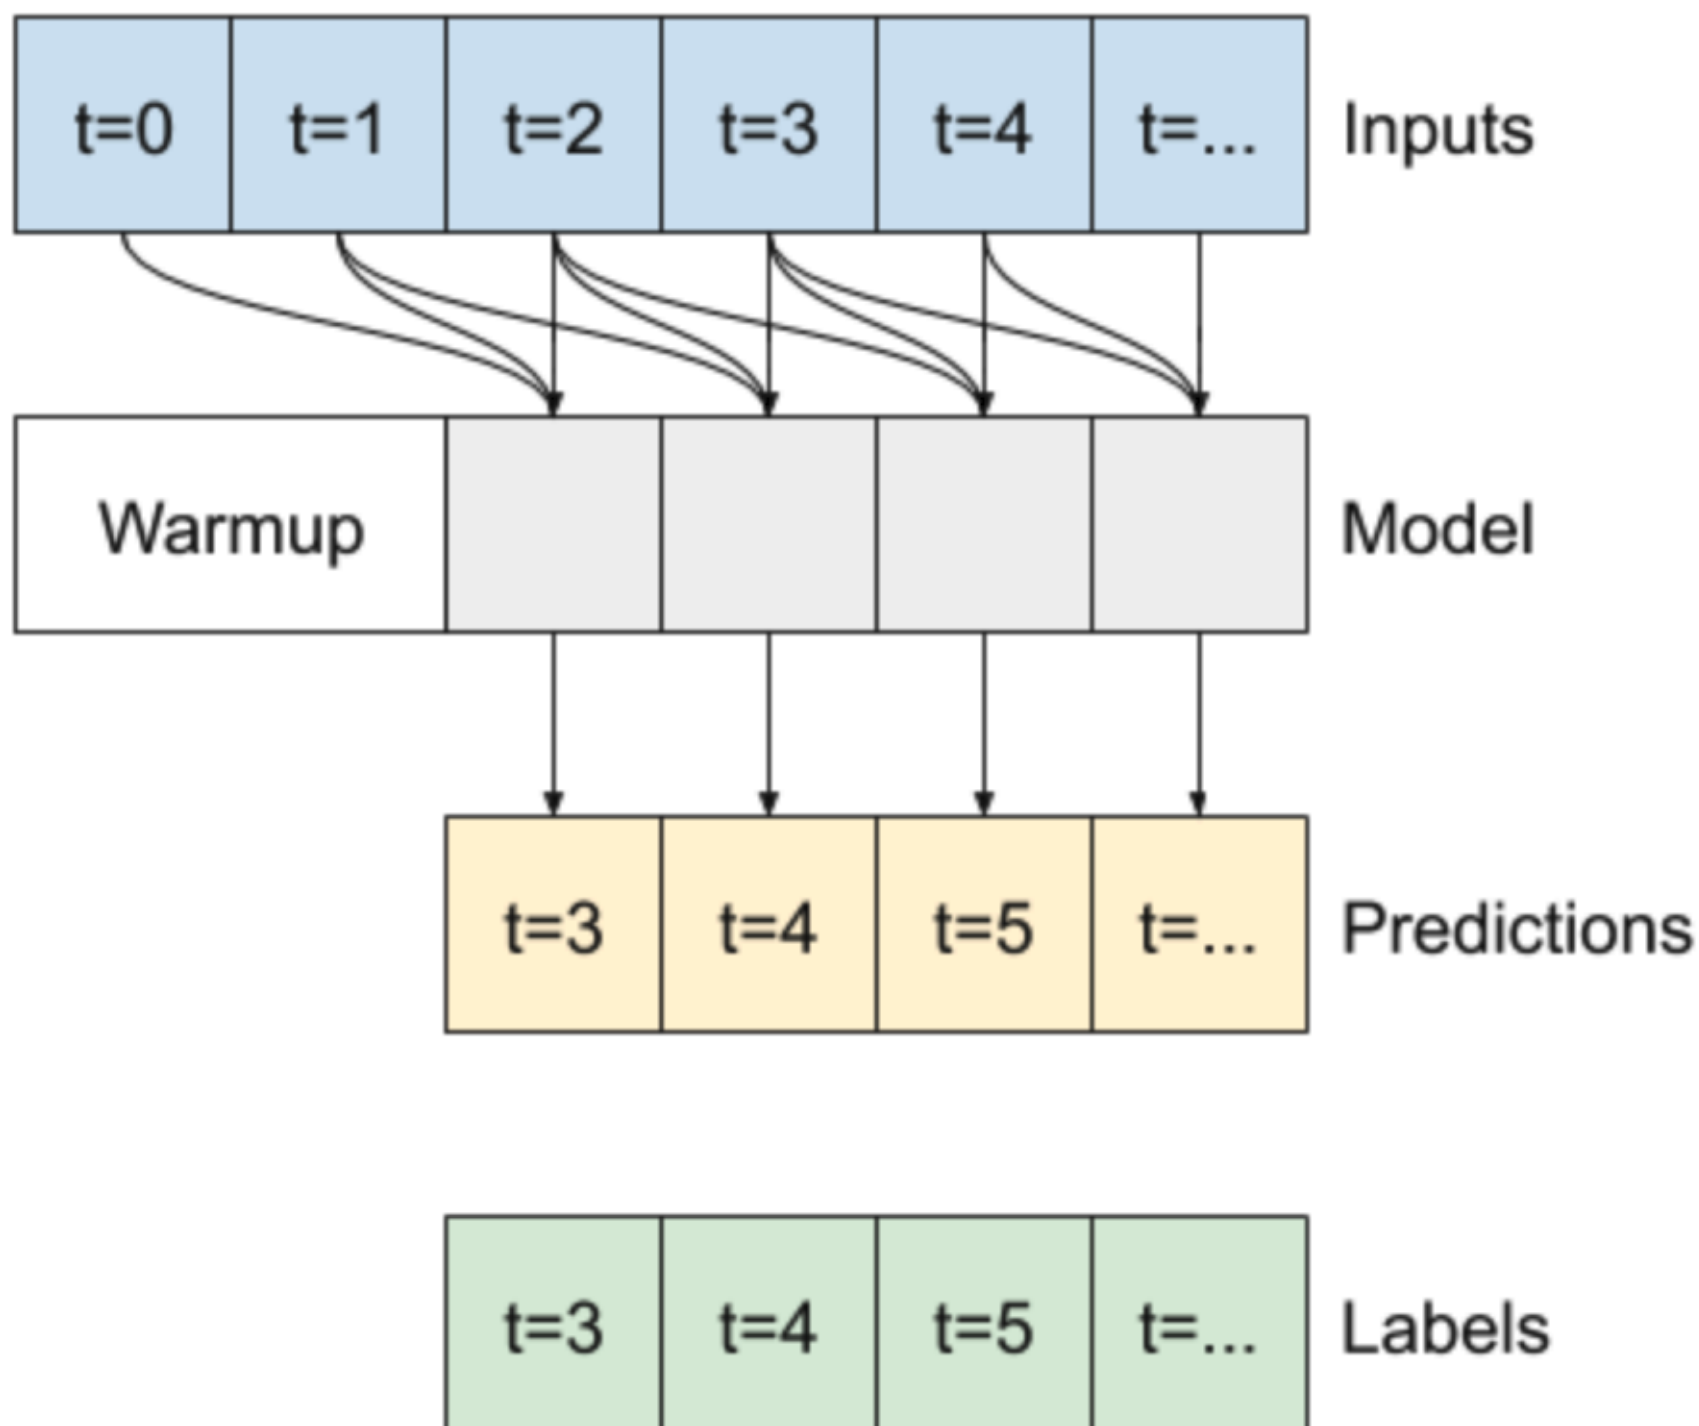
Correlation Matrix



# Expected Objective/Modelling

Time series modelling is not well modelled by traditional linear regression or decision trees due to the stochastic nature of the data. Because of this, we have to use modelling that has so-far not been used in this class. The current popular models for time series data are ARIMA (autoregressive integrated moving average) and LSTM Neural Networks (long short-term memory). We would like to put in a the features created in the last section and have an output of the total cost for the next month.





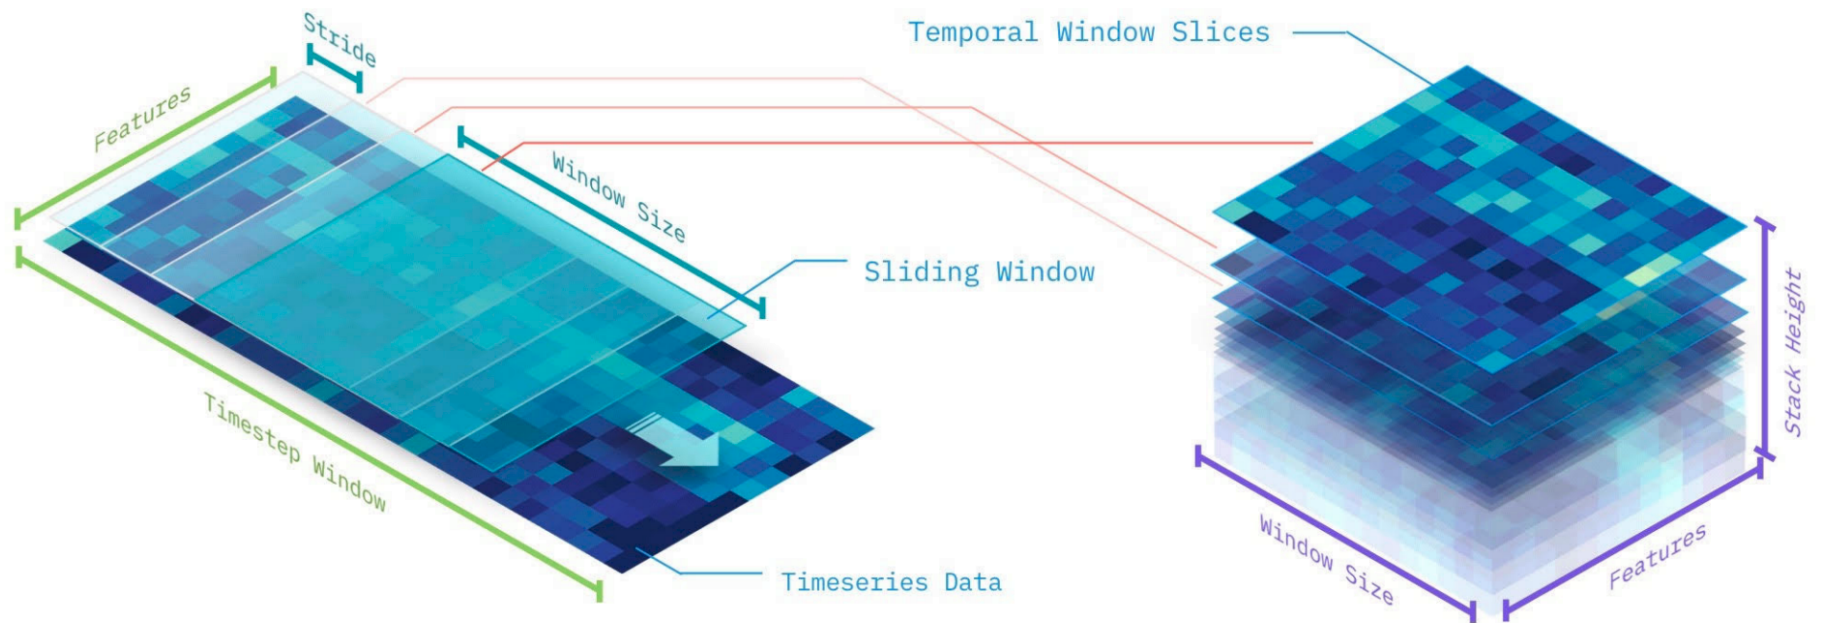
```
In [6]: def to_supervised (df_for_training, n_future, n_past):
trainX = []
trainY = []

for i in range(n_past, len(df_for_training) - n_future + 1):
    trainX.append(df_for_training[i - n_past:i, 0:df_for_training.shape[1]])
    trainY.append(df_for_training[i + n_future - 1:i + n_future, 0])

trainX, trainY = np.array(trainX), np.array(trainY)

print('trainX shape == {}'.format(trainX.shape))
print('trainY shape == {}'.format(trainY.shape))

return trainX, trainY
```



In [7]:

```
def build_model (df_for_training, n_future, n_past):

    trainX, trainY = to_supervised (df_for_training, n_future, n_past)

    model = Sequential()
    model.add(LSTM(64, activation='relu', input_shape=(trainX.shape[1], trainX.shape[2]), return_sequences=True))
    model.add(LSTM(32, activation='relu', return_sequences=False))
    model.add(Dropout(0.2))
    model.add(Dense(trainY.shape[1]))

    model.compile(optimizer='adam', loss='mse') #custom loss function, l2/l1 regularization
    model.summary()

    es = EarlyStopping(monitor='val_loss', min_delta=1e-10, patience=10, verbose=1)
    rlr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=10, verbose=1)
    mcp = ModelCheckpoint(filepath='weights.h5', monitor='val_loss', verbose=1, save_best_only=True, save_weights_only=True)

    tb = TensorBoard('logs')

    history = model.fit(trainX, trainY, shuffle=True, epochs=30, callbacks=[es, rlr, mcp, tb], validation_split=0.2, verbose=1)
    plt.plot(history.history['loss'], label='Training loss')
    plt.plot(history.history['val_loss'], label='Validation loss')
    plt.legend()

    return model, trainX, trainY
```